

Dynatrace Associate Certification Practice Question Bank

200 Original Multiple-Choice Questions with Answers & Explanations

Covering: Platform Fundamentals - OneAgent - ActiveGate - Infrastructure Monitoring
- Smartscape - Tagging & Management Zones - Davis AI & Problems - RUM -
Synthetic - Dashboards & DQL - Administration

StackMaster Tech • Corporate Training

Original practice material for exam preparation. Not affiliated with or endorsed by Dynatrace; these are not actual exam questions.

How to Use This Question Bank

Each question is followed immediately by the correct answer and a one-line explanation. For self-testing, cover the answer line, attempt the question, then verify. Questions are grouped into 11 sections mirroring the Dynatrace Associate exam blueprint areas. Aim for 85%+ accuracy per section before booking the exam, and revisit weak sections using Dynatrace documentation and hands-on labs.

Platform Fundamentals & Architecture (20 questions)

Q1. Which component of the Dynatrace platform is responsible for automatic root-cause analysis?

- A. Smartscape
- B. Davis AI**
- C. OneAgent
- D. ActiveGate

Answer: B — Davis is Dynatrace's causal AI engine that performs automatic root-cause analysis on detected problems.

Q2. What is Grail in the Dynatrace platform?

- A. A synthetic testing engine
- B. A unified data lakehouse for logs, metrics, traces, and events**
- C. A browser plugin
- D. A deployment automation tool

Answer: B — Grail is the massively parallel data lakehouse that stores observability data without pre-indexing.

Q3. Which deployment model does Dynatrace SaaS use?

- A. Customer-hosted cluster
- B. Dynatrace-hosted cloud environment**
- C. On-premises appliance only
- D. Agentless-only model

Answer: B — In SaaS, the Dynatrace cluster is hosted and operated by Dynatrace in the cloud.

Q4. Dynatrace Managed differs from SaaS primarily in that:

- A. It cannot monitor cloud workloads
- B. The cluster runs on customer-controlled infrastructure**
- C. It has no Davis AI
- D. It supports only Windows

Answer: B — Managed runs the Dynatrace cluster on infrastructure the customer controls, while Dynatrace still delivers updates.

Q5. Which language is used to query data stored in Grail?

- A. SQL only
- B. DQL (Dynatrace Query Language)**
- C. SPL
- D. PromQL only

Answer: B — DQL is the pipeline-based query language used to fetch and process data from Grail.

Q6. What does the term 'environment' mean in Dynatrace?

- A. A single monitored host
- B. A monitoring instance with its own data, settings, and users
- C. A Kubernetes namespace
- D. A dashboard folder

Answer: B — An environment is an isolated monitoring instance identified by its own environment ID and URL.

Q7. Which of the following is NOT a core data type ingested by the Dynatrace platform?

- A. Metrics
- B. Traces
- C. Logs
- D. Compiled binaries

Answer: D — Dynatrace ingests metrics, traces, logs, and events - not compiled application binaries.

Q8. PurePath technology in Dynatrace refers to:

- A. A network routing protocol
- B. End-to-end distributed trace capture at code level
- C. A dashboard template
- D. A billing model

Answer: B — PurePath captures end-to-end, code-level distributed transactions across tiers.

Q9. Which pillar of observability do PurePaths primarily represent?

- A. Metrics
- B. Logs
- C. Traces
- D. Events

Answer: C — PurePaths are Dynatrace's implementation of distributed traces.

Q10. The Dynatrace web UI is organized into 'apps' in the latest platform. Which app is used to query and explore any data with DQL?

- A. Notebooks
- B. Kubernetes app
- C. Clouds app
- D. Hub

Answer: A — Notebooks let users run DQL queries and combine results with text and visualizations.

Q11. What is the Dynatrace Hub used for?

- A. Managing user passwords
- B. Discovering and activating extensions, apps, and technology integrations
- C. Storing PurePaths
- D. Editing firewall rules

Answer: B — The Hub is the in-product catalog for extensions, apps, and integrations.

Q12. Which statement about Dynatrace licensing consumption is TRUE?

- A. All features consume the same unit
- B. Different capabilities (e.g., host monitoring, DEM, log ingest) consume different units
- C. Licensing is per dashboard
- D. Licensing is per user login

Answer: B — Consumption is tracked per capability, such as host units/hours, DEM units, and log ingest volume.

Q13. A 'host unit' in classic Dynatrace licensing is based on:

- A. Number of CPUs only
- B. RAM size of the monitored host
- C. Disk capacity
- D. Network throughput

Answer: B — Host unit consumption is calculated from the amount of RAM on the monitored host.

Q14. Which Dynatrace component provides the visual, real-time topology map of your environment?

- A. Smartscape
- B. Session Replay
- C. Hub
- D. Grail

Answer: A — Smartscape visualizes the vertical and horizontal dependencies of all monitored entities.

Q15. In the Dynatrace entity model, which of the following is a monitored entity type?

- A. HOST
- B. SERVICE
- C. PROCESS_GROUP_INSTANCE
- D. All of the above

Answer: D — Hosts, services, and process group instances are all first-class entity types.

Q16. What uniquely identifies every monitored entity in Dynatrace?

- A. Its display name
- B. Its entity ID (e.g., HOST-ABC123)
- C. Its IP address
- D. Its tag

Answer: B — Every entity has a unique entity ID; names and IPs can change or overlap.

Q17. Which protocol do OneAgent and ActiveGate use to send data to the Dynatrace cluster?

- A. FTP
- B. HTTPS (encrypted)
- C. SNMP
- D. Telnet

Answer: B — All communication to the cluster is encrypted over HTTPS.

Q18. What is the primary purpose of the Dynatrace API tokens?

- A. To style dashboards
- B. To authenticate programmatic access to Dynatrace APIs with scoped permissions
- C. To encrypt PurePaths
- D. To license hosts

Answer: B — API tokens carry scopes that authorize specific API operations.

Q19. Which of these is a valid reason to choose Dynatrace Managed over SaaS?

- A. Faster UI rendering
- B. Data residency and regulatory requirements to keep data on-premises
- C. Free licensing
- D. No need for updates

Answer: B — Managed is typically selected when data must remain within customer-controlled infrastructure.

Q20. The 'Environment ID' appears in which part of a SaaS URL?

- A. After the # fragment
- B. As the subdomain, e.g., abc12345.live.dynatrace.com
- C. In the query string only
- D. It never appears in URLs

Answer: B — SaaS environments are addressed via their environment ID as the subdomain of live.dynatrace.com.

OneAgent Deployment & Configuration (30 questions)

Q21. What is the primary function of OneAgent?

- A. To store monitoring data
- B. To automatically collect metrics, traces, and logs from a host and its processes
- C. To render dashboards
- D. To manage user permissions

Answer: B — OneAgent is installed on hosts to auto-instrument processes and collect full-stack telemetry.

Q22. Which installation modes does OneAgent support?

- A. Full-stack only
- B. Full-stack, Infrastructure, and Discovery modes
- C. Log-only mode
- D. Synthetic mode

Answer: B — OneAgent can run in Full-stack, Infrastructure-only, or Discovery mode depending on monitoring needs.

Q23. In Infrastructure monitoring mode, OneAgent does NOT collect:

- A. Host CPU metrics
- B. Process availability
- C. Code-level traces (PurePaths)
- D. Host network metrics

Answer: C — Infrastructure mode excludes deep code-level instrumentation such as PurePaths.

Q24. How does OneAgent instrument most application processes on a host?

- A. Manual code changes per app
- B. Automatic injection into supported process technologies
- C. Recompiling the application
- D. Byte-copying binaries to Dynatrace

Answer: B — OneAgent automatically injects code modules into supported technologies without code changes.

Q25. Which privilege level is required to install OneAgent on Linux?

- A. Any regular user
- B. root (or sudo)
- C. Only the dynatrace user
- D. No privileges needed

Answer: B — Installation requires root/sudo to deploy system-level components.

Q26. After installing OneAgent, deep monitoring of an already-running process typically requires:

- A. A cluster restart
- B. A restart of the monitored process
- C. Reinstalling the OS
- D. Nothing at all in every case

Answer: B — Injection into a process generally requires the process to be restarted after installation.

Q27. What is a 'network zone' used for in OneAgent configuration?

- A. Blocking internet access
- B. Routing OneAgent traffic to preferred ActiveGates
- C. Renaming hosts
- D. Encrypting disks

Answer: B — Network zones let OneAgents prefer specific ActiveGates for communication.

Q28. Which parameter assigns a custom host name during OneAgent installation?

- A. --set-host-name
- B. --set-network-zone
- C. --set-app-log-content-access
- D. --set-server

Answer: A — --set-host-name overrides the detected host name with a custom one.

Q29. Which OneAgent command-line tool is used to change configuration after installation?

- A. oneagentctl
- B. dtctl
- C. agentmgr
- D. dynactl

Answer: A — oneagentctl reads and changes OneAgent settings such as host name, tags, and endpoints.

Q30. How can host-level tags be added at install time?

- A. --set-host-tag=
- B. --add-label
- C. --tag-mode
- D. Tags cannot be set at install

Answer: A — The installer (and oneagentctl) supports --set-host-tag to add host tags.

Q31. Which statement about OneAgent updates is TRUE?

- A. Updates always require manual reinstall
- B. OneAgent can auto-update based on environment update settings
- C. Updates are impossible on Windows
- D. Updates require cluster downtime

Answer: B — Auto-update policies let OneAgents update themselves to target versions automatically.

Q32. Where do you verify that a host's OneAgent is connected and sending data?

- A. Deployment status page in Dynatrace
- B. The billing console
- C. Session Replay
- D. The Hub

Answer: A — Deployment status shows all OneAgents, versions, and connectivity state.

Q33. Which of the following technologies is auto-instrumented for deep monitoring by OneAgent?

- A. Java
- B. Static HTML files
- C. CSV files
- D. BIOS firmware

Answer: A — Java is one of the core technologies OneAgent instruments at code level.

Q34. What is the purpose of the OneAgent 'monitoring rules' for processes?

- A. To define which processes receive deep monitoring injection
- B. To create dashboards
- C. To rotate logs
- D. To set user passwords

Answer: A — Process monitoring rules control which processes are injected for deep monitoring.

Q35. A host shows 'OneAgent not running'. Which is the FIRST reasonable troubleshooting step?

- A. Reinstall the OS
- B. Check the OneAgent service/daemon status on the host
- C. Delete the environment
- D. Disable the firewall permanently

Answer: B — Verifying the local OneAgent service state is the standard first diagnostic step.

Q36. Which directory contains OneAgent logs on Linux by default?

- A. /opt/dynatrace/oneagent/log
- B. /var/lib/dynatrace/oneagent/log
- C. /var/log/dynatrace/oneagent
- D. /etc/dynatrace/logs

Answer: C — OneAgent writes its logs under /var/log/dynatrace/oneagent by default.

Q37. What does 'app-log content access' control?

- A. Whether OneAgent may read and ingest log file content from monitored processes
- B. Access to the Dynatrace UI
- C. API token scopes
- D. Dashboard sharing

Answer: A — This setting governs whether log content on the host can be accessed and ingested.

Q38. Which environments can OneAgent monitor?

- A. Physical hosts only
- B. VMs only
- C. Physical, virtual, cloud instances, and containers/Kubernetes nodes
- D. Mainframes only

Answer: C — OneAgent supports physical, virtual, cloud, and containerized infrastructure.

Q39. In Kubernetes, which component typically deploys OneAgent to every node?

- A. A single pod
- B. Dynatrace Operator (e.g., via DynaKube custom resource)
- C. kubectl cp
- D. A cron job

Answer: B — The Dynatrace Operator manages OneAgent rollout across cluster nodes via the DynaKube CR.

Q40. What is 'cloudNativeFullStack' in the DynaKube configuration?

- A. A billing tier
- B. A deployment mode combining node-level OneAgents with pod-level code modules
- C. A dashboard preset
- D. A log format

Answer: B — cloudNativeFullStack deploys infrastructure monitoring per node plus application-only injection per pod.

Q41. Which token is required when installing OneAgent so it can register with the environment?

- A. Personal access token
- B. PaaS/InstallerDownload + tenant token embedded in installer
- C. SSL certificate
- D. Kerberos ticket

Answer: B — The installer carries a tenant token that authenticates the agent to the environment.

Q42. What happens if OneAgent cannot reach the Dynatrace cluster temporarily?

- A. Data is lost immediately
- B. It buffers data locally and retries sending
- C. The host shuts down
- D. It uninstalls itself

Answer: B — OneAgent buffers and retries, tolerating short network outages.

Q43. Which OS families are supported by OneAgent?

- A. Windows only
- B. Linux only
- C. Windows, Linux, AIX, and more
- D. DOS

Answer: C — OneAgent supports Windows, many Linux distributions, AIX, and other platforms.

Q44. The 'Discovery' mode of OneAgent is primarily intended to:

- A. Provide full APM tracing
- B. Give lightweight visibility for identifying monitoring candidates at minimal cost
- C. Replace ActiveGate
- D. Capture session replays

Answer: B — Discovery mode offers lightweight insights to decide where to enable deeper monitoring.

Q45. Which oneagentctl command shows the currently configured communication endpoints?

- A. oneagentctl --get-server
- B. oneagentctl --show-net
- C. oneagentctl --endpoints
- D. oneagentctl --get-url

Answer: A — --get-server prints the configured Dynatrace communication endpoints.

Q46. Host groups in OneAgent deployment are used to:

- A. Group hosts for shared configuration such as process monitoring and naming rules
- B. Encrypt traffic
- C. Create users
- D. Store dashboards

Answer: A — Host groups let you apply settings consistently to sets of hosts.

Q47. How do you place a host into a host group?

- A. Only via the UI after install
- B. Installer/oneagentctl parameter --set-host-group
- C. Editing Smartscape
- D. It is automatic and cannot be changed

Answer: B — Host group membership is set at install time or via oneagentctl --set-host-group (restart required).

Q48. Which of the following does OneAgent capture WITHOUT any configuration for supported web servers?

- A. Incoming request traces and response times
- B. Employee HR records
- C. Database table contents
- D. Source code repositories

Answer: A — Request-level tracing for supported servers works out of the box.

Q49. A process is supported but shows no code-level visibility. A common cause is:

- A. The process started before OneAgent installation and was not restarted
- B. Too many dashboards
- C. Wrong timezone
- D. Excessive tagging

Answer: A — Processes running before installation need a restart for injection.

Q50. Which statement about OneAgent resource consumption is accurate?

- A. It is designed for low overhead in production use
- B. It requires 50% of host CPU
- C. It disables host swapping
- D. It duplicates all application memory

Answer: A — OneAgent is engineered for minimal production overhead.

ActiveGate (15 questions)

Q51. What is the primary role of an Environment ActiveGate?

- A. Replace OneAgent
- B. Act as a secure proxy that routes, compresses, and encrypts OneAgent traffic to the cluster
- C. Render the UI
- D. Store long-term data

Answer: B — Environment ActiveGate proxies and optimizes agent-to-cluster communication.

Q52. Which capability requires an ActiveGate rather than only OneAgent?

- A. Host CPU monitoring
- B. Private synthetic monitoring locations
- C. Process injection
- D. Log writing

Answer: B — Private synthetic locations run on ActiveGates with the synthetic capability enabled.

Q53. Cloud integrations such as Azure or VMware monitoring without agents on every VM are executed by:

- A. Session Replay
- B. ActiveGate remote monitoring extensions
- C. Smartscape
- D. Notebooks

Answer: B — ActiveGate hosts extensions that poll cloud/virtualization APIs remotely.

Q54. Which statement about ActiveGate placement is a best practice?

- A. Place it outside all firewalls with public access to everything
- B. Deploy it inside the network so many OneAgents share one secured egress point
- C. Install one per monitored process
- D. Only install on laptops

Answer: B — A few ActiveGates inside the network consolidate and secure outbound traffic.

Q55. What is a Cluster ActiveGate (Dynatrace Managed)?

- A. A browser extension
- B. An ActiveGate that provides access to the cluster, e.g., for agents outside the network or mobile RUM
- C. A database
- D. A CI/CD tool

Answer: B — Cluster ActiveGates front the Managed cluster for external traffic scenarios.

Q56. Which port does ActiveGate listen on by default for OneAgent traffic?

- A. 22
- B. 443 (HTTPS) / 9999 internal
- C. 80 only
- D. 3306

Answer: B — ActiveGate communication is HTTPS-based; 443 is the standard external port (9999 used in some internal setups).

Q57. Kubernetes API monitoring (via the Kubernetes integration) typically runs on:

- A. A containerized ActiveGate inside the cluster
- B. The developer laptop
- C. The etcd pod itself
- D. A mainframe

Answer: A — A containerized ActiveGate (or Operator-managed component) queries the Kubernetes API.

Q58. Which is NOT an ActiveGate capability/module?

- A. Synthetic monitoring
- B. Metric/API ingestion endpoint
- C. z/OS hub
- D. Compiling application source code

Answer: D — ActiveGate has modules like synthetic and API endpoints; it never compiles customer code.

Q59. If OneAgents must not access the internet directly, the recommended pattern is:

- A. Disable monitoring
- B. Route OneAgent traffic through an Environment ActiveGate
- C. Use FTP
- D. Email the data

Answer: B — ActiveGate serves as the single controlled egress for agent traffic.

Q60. Memory dumps and other on-demand diagnostics from hosts are routed through:

- A. Smartscape
- B. ActiveGate
- C. Grail directly
- D. The Hub

Answer: B — Diagnostic data such as memory dumps flows via ActiveGate.

Q61. How is an ActiveGate installed?

- A. Only as a Docker image
- B. Via an installer downloaded from Deployment status > ActiveGates (Windows/Linux) or as a container
- C. Via the mobile app
- D. It is built into every OS

Answer: B — ActiveGate installers are provided in the deployment pages, plus container options.

Q62. Multiple ActiveGates in the same network zone provide:

- A. Nothing extra
- B. Load balancing and failover for OneAgent traffic
- C. Slower routing
- D. Duplicate data

Answer: B — OneAgents can fail over among ActiveGates within a zone.

Q63. Which component exposes a local API ingest endpoint for metrics/logs inside a private network?

- A. OneAgent GUI
- B. ActiveGate (API endpoint module)
- C. Session Replay
- D. Smartscape

Answer: B — ActiveGate can receive API ingest locally and forward it to the cluster.

Q64. An ActiveGate version should generally be:

- A. Older than OneAgents
- B. Equal to or newer than the OneAgents that connect through it
- C. Irrelevant
- D. Always beta

Answer: B — Keeping ActiveGate at least as new as connecting agents avoids compatibility issues.

Q65. Where do you check ActiveGate connectivity and version in the UI?

- A. Deployment status > ActiveGates
- B. Session segmentation
- C. Service flow
- D. Hub billing tab

Answer: A — Deployment status lists all ActiveGates with health and version details.

Infrastructure Monitoring: Hosts, Processes & Services (25 questions)

Q66. Which metrics are captured for every monitored host by default?

- A. CPU, memory, disk, and network
- B. Only CPU
- C. Only disk
- D. Only uptime

Answer: A — Host monitoring includes CPU, memory, disk, NIC, and availability out of the box.

Q67. A 'process group' in Dynatrace is:

- A. A random set of processes
- B. Logically equivalent processes (same technology/purpose) merged across instances
- C. A Linux cgroup
- D. A user group

Answer: B — Process groups cluster equivalent processes, e.g., all instances of the same app server.

Q68. What is the difference between a process group and a process group instance (PGI)?

- A. No difference
- B. PGI is one concrete running process on a host; the group aggregates them
- C. PGI is a dashboard
- D. Group is per-thread

Answer: B — A PGI is a specific instance on a host; the process group spans all such instances.

Q69. Services in Dynatrace are detected primarily from:

- A. Manual entry
- B. Instrumented server-side entry points (e.g., web requests, queues, DB calls)
- C. Ping tests
- D. DNS records

Answer: B — Services are derived from monitored process entry points like web request handlers.

Q70. Which built-in service metrics are provided automatically?

- A. Response time, failure rate, throughput
- B. Only license usage
- C. Only disk IO
- D. Employee count

Answer: A — Golden signals such as response time, failure rate, and request count come out of the box.

Q71. Host availability states include:

- A. Only up
- B. Monitored/up, unmonitored, shut down (and related states)
- C. Blue/green
- D. Hot/cold

Answer: B — Dynatrace distinguishes running/monitored, shutdown, and unmonitored states.

Q72. Process availability monitoring policies allow you to:

- A. Alert when specified critical processes are not running on a host
- B. Encrypt processes
- C. Speed up processes
- D. Rename hosts

Answer: A — Availability policies detect and alert on missing critical processes by detection rules.

Q73. Which technology detail does Dynatrace show for a process?

- A. Technology type and version (e.g., Java 17, Tomcat)
- B. Purchase price
- C. Developer names
- D. Git history

Answer: A — Process pages display detected technologies and versions.

Q74. Disk-related default alerts include:

- A. Low disk space and inode exhaustion
- B. Disk color
- C. Disk vendor
- D. Disk temperature only

Answer: A — Built-in events cover low disk space and low inodes among others.

Q75. What does 'CPU saturation' on a host page indicate?

- A. CPU demand exceeds capacity, often with high load/queue
- B. CPU is idle
- C. CPU is disabled
- D. A billing state

Answer: A — Saturation reflects contention where demand exceeds available CPU.

Q76. Network metrics per host include:

- A. Traffic in/out, retransmissions, connectivity quality
- B. Wi-Fi passwords
- C. Router firmware
- D. Cable length

Answer: A — Host networking shows throughput and TCP quality indicators like retransmissions.

Q77. Log Monitoring in Dynatrace can:

- A. Ingest, store in Grail, and query logs with DQL, linked to topology
- B. Only tail files locally
- C. Print logs
- D. Delete application logs

Answer: A — Logs are ingested to Grail, queryable via DQL, and connected to entities.

Q78. Which statement about custom metrics ingestion is TRUE?

- A. Impossible
- B. Supported via the Metrics API / OneAgent metric ingestion (e.g., dynatrace_ingest, StatsD-like line protocol)
- C. Only via email
- D. Only for Windows

Answer: B — Custom metrics can be pushed via API or local ingestion interfaces.

Q79. The 'Multidimensional analysis' views for services let you:

- A. Slice response time/failures by dimensions like endpoint, or custom dimensions
- B. Change user passwords
- C. Deploy code
- D. Manage licenses

Answer: A — Multidimensional analysis breaks service metrics down by selected dimensions.

Q80. Which entity does a database call from an instrumented app create in Dynatrace?

- A. A database service (called service)
- B. A host only
- C. A dashboard
- D. A synthetic monitor

Answer: A — Database calls appear as database services with their own metrics.

Q81. VMware monitoring without installing agents on ESXi is achieved by:

- A. OneAgent on ESXi
- B. ActiveGate polling vCenter APIs
- C. SSH scripts
- D. Impossible

Answer: B — The VMware integration runs on ActiveGate and polls vCenter.

Q82. Which cloud integrations enrich monitoring with metadata and metrics via APIs?

- A. AWS, Azure, GCP integrations
- B. Only FTP
- C. Fax integration
- D. BIOS integration

Answer: A — Native cloud integrations pull metrics/metadata from AWS, Azure, and GCP.

Q83. What does 'full-stack' monitoring of a host add over infrastructure mode?

- A. Code-level service/PurePath visibility and deep process insights
- B. Just a different icon
- C. Lower data quality
- D. Nothing

Answer: A — Full-stack unlocks deep application tracing beyond infrastructure metrics.

Q84. Container monitoring with OneAgent provides:

- A. Per-container resource metrics and injection into supported containerized apps
- B. Only image names
- C. Docker Hub credentials
- D. Registry cleanup

Answer: A — OneAgent monitors container resources and instruments supported processes inside.

Q85. Which Kubernetes objects does the Dynatrace Kubernetes app surface?

- A. Clusters, nodes, namespaces, workloads, pods
- B. Only pods
- C. Only secrets
- D. Git branches

Answer: A — The Kubernetes experience covers cluster-to-pod level objects.

Q86. An unexpectedly high number of processes in one process group may be tuned via:

- A. Process group detection rules
- B. Deleting hosts
- C. Buying licenses
- D. Session Replay

Answer: A — Detection rules let you split or merge process groups deliberately.

Q87. Which of the following can trigger a custom event for alerting from a metric?

- A. Metric events (threshold/baseline based) configuration
- B. Smartscape click
- C. PDF export
- D. Renaming a dashboard

Answer: A — Metric events raise events/alerts when metric conditions are met.

Q88. 'Host anomalies' detection settings allow you to:

- A. Tune CPU/memory/disk anomaly detection thresholds per environment or host
- B. Overclock CPUs
- C. Delete anomalies
- D. Hide hosts

Answer: A — Anomaly detection for infrastructure can be tuned globally or per entity.

Q89. Which page shows all monitored technologies across processes in the environment?

- A. Technologies & processes overview
- B. Billing page
- C. User sessions
- D. Geolocation settings

Answer: A — The technology/processes overview aggregates detected technologies.

Q90. SNMP-based devices (e.g., network gear) are monitored using:

- A. OneAgent injection
- B. Extensions running on ActiveGate (Extension Framework)
- C. Session Replay
- D. Browser plugin

Answer: B — SNMP polling is done by extensions executed on ActiveGate.

Smartscape & Topology (10 questions)

Q91. Smartscape displays which two dependency directions?

- A. Vertical (stack: data center to application) and horizontal (calls between entities)
- B. North and south only
- C. Left and right only
- D. Alphabetical order

Answer: A — Smartscape shows vertical full-stack layers and horizontal call relationships.

Q92. Which layers appear in Smartscape?

- A. Applications, Services, Processes, Hosts, Data centers
- B. Only hosts
- C. Only users
- D. DNS zones

Answer: A — The classic layers span data centers up to applications.

Q93. How are Smartscape relationships discovered?

- A. Manually drawn
- B. Automatically from OneAgent-observed communications and topology data
- C. Imported from Visio
- D. Random

Answer: B — Topology is auto-detected from real observed traffic and system data.

Q94. Clicking a service in Smartscape shows:

- A. Its callers, callees, and underlying processes/hosts
- B. Its purchase order
- C. Its Git repo
- D. Nothing

Answer: A — Selection highlights inbound/outbound dependencies and the supporting stack.

Q95. Which analysis view traces the sequence of service-to-service calls for transactions?

- A. Service flow
- B. Disk view
- C. License view
- D. Geomap only

Answer: A — Service flow visualizes call sequences with volumes and response times.

Q96. The 'backtrace' of a service answers:

- A. Which callers/paths lead into this service
- B. Disk usage
- C. User passwords
- D. Billing

Answer: A — Backtrace shows who calls the service, back to applications.

Q97. Topology data in the newer platform is queryable via:

- A. Smartscape screenshots only
- B. Entity model / DQL and the Entities API
- C. FTP
- D. Email

Answer: B — Entities and relationships are accessible via APIs and DQL.

Q98. Which statement about Smartscape freshness is TRUE?

- A. It updates in near real time as topology changes
- B. It updates yearly
- C. It is static after install
- D. It needs manual refresh scripts

Answer: A — Smartscape continuously reflects current topology.

Q99. A host visible in Smartscape but with no services above it most likely runs:

- A. Infrastructure-only monitoring or no supported deep-monitored processes
- B. A full APM stack
- C. Synthetic monitors
- D. Session Replay

Answer: A — Without deep monitoring/entry points, no service layer appears above the host.

Q100. Why is automatic topology important for Davis AI?

- A. It provides the dependency context used to localize root cause
- B. It makes the UI colorful
- C. It is not used
- D. It stores passwords

Answer: A — Davis correlates problems along real dependencies to find root cause.

Tagging, Metadata & Management Zones (15 questions)

Q101. What is the main purpose of tags in Dynatrace?

- A. Organize, filter, and target entities for alerting, permissions, and analysis
- B. Encrypt data
- C. Speed up CPUs
- D. License hosts

Answer: A — Tags classify entities for filtering, alerting profiles, zones, and automation.

Q102. Automatically applied tags are created via:

- A. Auto-tagging rules based on entity properties/metadata
- B. Random assignment
- C. Only manual clicks
- D. Emails

Answer: A — Auto-tagging rules evaluate metadata (names, technology, cloud tags) to apply tags.

Q103. Which metadata sources can drive auto-tagging?

- A. Host names, process properties, Kubernetes labels, cloud tags, environment variables
- B. Only IP addresses
- C. Only usernames
- D. Fonts

Answer: A — Rules can use many detected properties, including Kubernetes labels and cloud tags.

Q104. A tag in Dynatrace can have the format:

- A. key or key:value
- B. Only numeric
- C. Only uppercase
- D. XML only

Answer: A — Tags may be simple keys or key:value pairs, optionally with contexts.

Q105. Management zones are used to:

- A. Partition the environment so teams see only relevant entities and data
- B. Delete entities
- C. Set CPU limits
- D. Change time zones

Answer: A — Zones scope visibility/filtering of entities, dashboards, and problems per team/app.

Q106. Management zone membership is defined by:

- A. Rules matching entities (tags, names, technologies, relationships)
- B. Physical location only
- C. Alphabet
- D. Payment tier

Answer: A — Zone rules select entities by properties, often via tags.

Q107. Which statement about management zones and permissions is TRUE?

- A. Zones can be bound to user groups to restrict what users can view
- B. Zones grant OS root
- C. Zones replace tokens
- D. Zones are per browser

Answer: A — Zone-based permissions limit user visibility to matching entities.

Q108. How do environment variables like DT_TAGS affect monitoring?

- A. They can add tags to processes/hosts automatically at startup
- B. They disable OneAgent
- C. They create users
- D. They format dashboards

Answer: A — DT_TAGS and similar variables inject tags/metadata from the deployment side.

Q109. DT_CUSTOM_PROP is used to:

- A. Attach custom key/value metadata properties to processes
- B. Set license keys
- C. Change UI language
- D. Rotate logs

Answer: A — DT_CUSTOM_PROP adds custom metadata usable in rules and tagging.

Q110. Why tag by 'ownership' or 'team'?

- A. To route alerts and scope zones/dashboards to responsible teams
- B. For fun
- C. To reduce RAM
- D. To disable Davis

Answer: A — Ownership tags drive alert routing, zoning, and accountability.

Q111. A single entity can belong to how many management zones?

- A. Exactly one
- B. Multiple zones simultaneously
- C. Zero always
- D. One per day

Answer: B — Entities may match several zone rules and appear in each.

Q112. Which is a good naming practice for tags at scale?

- A. Consistent key:value conventions, e.g., env:prod, app:checkout
- B. Random emojis
- C. One tag per employee
- D. No tags

Answer: A — Consistent structured tag conventions keep large environments manageable.

Q113. Manual tags differ from auto-tags in that:

- A. Manual tags are applied per entity by users; auto-tags come from rules and update dynamically
- B. Manual tags are faster CPUs
- C. Auto-tags are only for hosts
- D. No difference

Answer: A — Rule-based tags follow metadata changes; manual tags are static per entity.

Q114. Filtering a dashboard to one application team is best achieved with:

- A. A management zone or tag filter on the dashboard
- B. Deleting other data
- C. New environment per team always
- D. Screenshots

Answer: A — Zones/tag filters scope dashboard tiles to the relevant subset.

Q115. Kubernetes labels/annotations can propagate into Dynatrace as:

- A. Metadata usable for tagging and zone rules
- B. Nothing
- C. Only logs
- D. Billing items

Answer: A — K8s labels become entity metadata for rules and filtering.

Davis AI, Problems & Root Cause (20 questions)

Q116. Davis AI evaluates anomalies primarily against:

- A. Static values only
- B. Automatically learned baselines plus configured thresholds
- C. Coin flips
- D. Manual tickets

Answer: B — Davis combines auto-adaptive baselines with any static thresholds you define.

Q117. A Dynatrace 'problem' is best described as:

- A. A single raw alert
- B. A correlated group of related events/anomalies with impact and root-cause analysis
- C. A log line
- D. A dashboard

Answer: B — Problems aggregate causally related events instead of flooding separate alerts.

Q118. Which impact levels can a problem indicate?

- A. Infrastructure, Services, Applications (and real users/SLIs affected)
- B. Only color codes
- C. Font sizes
- D. None

Answer: A — Problems state affected layers and user/service impact.

Q119. The problem 'root cause' section shows:

- A. The entity/event Davis identified as the origin, with evidence
- B. Random entity
- C. Only time
- D. The bill

Answer: A — Davis pinpoints the likely root-cause entity and supporting evidence.

Q120. Which data does Davis correlate for root-cause analysis?

- A. Topology, events, metrics, traces, and logs in context
- B. Only host names
- C. Only tags
- D. Screenshots

Answer: A — Causal analysis spans dependencies and all telemetry in context.

Q121. Frequent-issue detection means:

- A. Recurring anomalies can be learned as 'frequent' and reduce alert noise
- B. Issues are emailed daily
- C. Problems duplicate
- D. Baselines reset hourly

Answer: A — Frequent issues are recognized to avoid repeated low-value alerts.

Q122. Availability, error, slowdown, and resource are examples of:

- A. Problem categories/types
- B. License tiers
- C. Browser types
- D. Zone names

Answer: A — Problems are classified into such categories based on the triggering events.

Q123. Where do you configure sensitivity for service response-time anomaly detection?

- A. Settings > Anomaly detection > Services
- B. BIOS
- C. Smartscape
- D. The Hub billing tab

Answer: A — Service anomaly detection settings tune slowdown/failure detection.

Q124. Custom threshold metric events differ from Davis baselines because they:

- A. Fire on user-defined static or seasonal conditions rather than learned baselines
- B. Cannot alert
- C. Only work at night
- D. Delete metrics

Answer: A — Metric events implement explicit conditions you define.

Q125. What is the purpose of 'alerting profiles'?

- A. Filter which problems trigger notifications, by severity, zone, tags, and delay
- B. Create users
- C. Encrypt alerts
- D. Rename problems

Answer: A — Alerting profiles select problem subsets for each notification channel.

Q126. Problem notifications can be sent via:

- A. Email, webhooks, Slack/Teams/PagerDuty/ServiceNow and other integrations
- B. Fax only
- C. Post
- D. Only UI

Answer: A — Numerous notification integrations plus generic webhooks are supported.

Q127. Which feature executes automated remediation when a problem opens?

- A. Workflows (AutomationEngine) or problem-notification-triggered automation
- B. Session Replay
- C. Geomaps
- D. Manual only

Answer: A — Workflows can trigger on problem events to run remediation actions.

Q128. Davis 'events' differ from 'problems' in that:

- A. Events are individual state changes/anomalies; problems are correlated event groups
- B. Events are bigger
- C. Problems come first
- D. No difference

Answer: A — Events feed the correlation that produces problems.

Q129. Maintenance windows are used to:

- A. Suppress alerting/problem detection during planned activities
- B. Reboot hosts
- C. Update browsers
- D. Bill customers

Answer: A — They prevent noise from expected maintenance impact.

Q130. The 'Problem evolution' view helps you:

- A. See how a problem developed over time across entities
- B. Change fonts
- C. Export licenses
- D. Create tokens

Answer: A — It replays the sequence of events and affected components.

Q131. Which statement about baselining for new services is TRUE?

- A. Davis needs a learning period before reliable baselines exist
- B. Baselines are instant and perfect
- C. Baselines never change
- D. Baselines are manual only

Answer: A — Auto-baselines require observation time to learn normal behavior.

Q132. An 'availability' problem for a host is raised when:

- A. The host or its monitoring becomes unavailable/shut down unexpectedly
- B. CPU is 10%
- C. A tag is added
- D. A user logs in

Answer: A — Unexpected unavailability triggers availability problems.

Q133. Root-cause relevance of log data means:

- A. Related log anomalies/errors can appear as evidence within a problem
- B. Logs are deleted
- C. Logs replace metrics
- D. Logs disable Davis

Answer: A — Logs in context strengthen root-cause evidence.

Q134. How can you reduce duplicate/noisy alerts across teams?

- A. Use correlation (problems), alerting profiles, zones, and frequent-issue detection
- B. Send everything to everyone
- C. Disable monitoring
- D. More emails

Answer: A — Scoped profiles plus Davis correlation cut noise dramatically.

Q135. The Davis 'why' explanation supports which operational goal?

- A. Faster MTTR through explainable, evidence-backed root cause
- B. Slower triage
- C. Hidden causes
- D. More tickets

Answer: A — Explainability accelerates diagnosis and repair.

Real User Monitoring (RUM) (15 questions)

Q136. Real User Monitoring captures:

- A. Actual end-user interactions, timings, and errors in web/mobile apps
- B. Only server CPU
- C. Synthetic scripts
- D. Firewall rules

Answer: A — RUM measures real user sessions, actions, and experience.

Q137. For web applications served by OneAgent-monitored web servers, the RUM JavaScript is:

- A. Injected automatically into pages
- B. Typed by users
- C. Mailed as a file
- D. Not needed

Answer: A — Auto-injection inserts the RUM snippet without code changes.

Q138. If auto-injection is not possible, RUM can be enabled via:

- A. Manual insertion of the JavaScript snippet/agentless RUM
- B. Only synthetic
- C. SNMP
- D. BIOS

Answer: A — Agentless/manual snippet setup covers non-injected scenarios.

Q139. A 'user action' in RUM is:

- A. An interaction like a page load, click/XHR that triggers activity
- B. A support ticket
- C. A CPU cycle
- D. A tag

Answer: A — User actions are load/interaction events measured end to end.

Q140. A 'user session' is:

- A. A sequence of user actions from one user visit, ending after inactivity or limits
- B. One CPU second
- C. An API token
- D. A log file

Answer: A — Sessions group a visitor's consecutive actions.

Q141. Apdex in Dynatrace rates:

- A. User experience of actions as satisfied/tolerating/frustrated
- B. Disk speed
- C. License usage
- D. Tag quality

Answer: A — Apdex classifies action performance into satisfaction bands.

Q142. Key user actions allow you to:

- A. Track/alert and set custom Apdex on business-critical actions
- B. Delete actions
- C. Hide users
- D. Change DNS

Answer: A — Marking key actions elevates monitoring for critical journeys.

Q143. Session Replay provides:

- A. Video-like playback of real user sessions with privacy controls
- B. Live webcam
- C. Screen sharing
- D. Server console

Answer: A — Replay reconstructs user sessions for UX/error analysis with masking options.

Q144. Which privacy capability is essential in Session Replay?

- A. Masking/blocking sensitive data (input fields, elements)
- B. Publishing data
- C. Disabling HTTPS
- D. IP spoofing

Answer: A — Masking rules prevent capture of sensitive user content.

Q145. RUM for mobile apps requires:

- A. Instrumenting the app with OneAgent SDK/mobile agent
- B. Nothing
- C. Synthetic only
- D. ActiveGate on the phone

Answer: A — Mobile RUM uses in-app instrumentation.

Q146. Which errors does RUM detect on the client side?

- A. JavaScript errors and request/HTTP errors affecting users
- B. Kernel panics
- C. Printer errors
- D. BIOS errors

Answer: A — Client-side script and request errors are captured per action/session.

Q147. Geolocation breakdown in RUM helps to:

- A. Analyze performance/experience by user region
- B. Set time zones
- C. Encrypt data
- D. Pick fonts

Answer: A — Geo views localize experience issues by region/ISP.

Q148. User session properties can be enriched with:

- A. Custom session/action properties (e.g., customer tier) for segmentation
- B. CPU registers
- C. BIOS strings
- D. Nothing

Answer: A — Custom properties enable business-centric segmentation.

Q149. Conversion goals in RUM measure:

- A. Completion of defined user journeys/destinations
- B. Server patching
- C. Backups
- D. Token expiry

Answer: A — Goals track business outcomes like reaching checkout success.

Q150. How does RUM connect to server-side traces?

- A. User actions link to backend PurePaths for full-journey analysis
- B. It does not
- C. Via email
- D. Only in PDFs

Answer: A — End-to-end linking ties browser actions to server-side traces.

Synthetic Monitoring (15 questions)

Q151. Synthetic monitoring differs from RUM in that it:

- A. Proactively simulates user or API traffic on a schedule, independent of real users
- B. Records real users
- C. Requires end users online
- D. Is only for mobile

Answer: A — Synthetic runs scripted checks even with zero real traffic - ideal for availability SLAs.

Q152. Which synthetic monitor type performs a scripted multi-step browser journey?

- A. Browser clickpath monitor
- B. Ping
- C. SNMP trap
- D. Log monitor

Answer: A — Clickpaths replay recorded browser interactions step by step.

Q153. A single-URL browser monitor checks:

- A. Load/availability/performance of one page in a real browser
- B. A TCP port only
- C. An email inbox
- D. Host CPU

Answer: A — It opens one URL in a browser and measures availability and timing.

Q154. HTTP monitors are best suited for:

- A. API endpoint availability and response validation without a full browser
- B. Video streaming tests
- C. Session Replay
- D. Disk checks

Answer: A — HTTP monitors issue lightweight requests, ideal for APIs and health endpoints.

Q155. Where can synthetic monitors execute from?

- A. Public Dynatrace locations and private locations (on ActiveGate)
- B. Only your laptop
- C. Only US
- D. Only via VPN apps

Answer: A — Both global public locations and private ActiveGate-based locations are supported.

Q156. Private synthetic locations are required when:

- A. Testing internal apps not reachable from the public internet
- B. Testing google.com
- C. Never
- D. Only for mobile

Answer: A — Internal endpoints require executing from inside the network.

Q157. Synthetic availability is typically calculated as:

- A. Percentage of successful executions over total executions in a period
- B. CPU/RAM ratio
- C. Ping latency
- D. License units

Answer: A — Availability derives from successful vs. failed monitor runs.

Q158. Which validation can an HTTP monitor perform?

- A. Response code, response body content/pattern, and performance thresholds
- B. Font checks
- C. Disk defragmentation
- D. BIOS versions

Answer: A — HTTP monitors validate status codes, content, and timing.

Q159. A clickpath step fails on a changed button locator. Best practice is to:

- A. Update/re-record the affected step or use resilient locators
- B. Delete synthetic entirely
- C. Ignore forever
- D. Reboot ActiveGate

Answer: A — Maintaining locators/scripts keeps clickpaths reliable through UI changes.

Q160. Synthetic monitor frequency affects:

- A. Detection speed of outages and consumption of synthetic units
- B. Nothing
- C. Host RAM tags
- D. Time zones

Answer: A — Higher frequency detects issues faster but consumes more units.

Q161. Which outage handling option avoids false alarms from a single flaky run?

- A. Retry on error / require consecutive failures before problem creation
- B. Alert on first ping
- C. Disable alerts
- D. Email everyone

Answer: A — Retries and consecutive-failure logic reduce false positives.

Q162. Performance thresholds on synthetic monitors can:

- A. Raise problems when steps exceed defined durations
- B. Change DNS
- C. Resize VMs
- D. Order pizza

Answer: A — Threshold breaches generate performance events/problems.

Q163. Comparing synthetic results across locations helps identify:

- A. Regional/network-specific issues vs. global outages
- B. License gaps
- C. Tag typos
- D. User names

Answer: A — Multi-location data separates local network problems from app-wide failures.

Q164. Credential vault in synthetic settings is used to:

- A. Securely store credentials/certificates used by monitors
- B. Store user photos
- C. Hold API bills
- D. Encrypt disks

Answer: A — The vault manages secrets referenced by synthetic scripts safely.

Q165. Which is a typical use case pairing synthetic with RUM?

- A. Baseline 24x7 availability via synthetic plus real-experience insight via RUM
- B. Replacing all monitoring with synthetic
- C. Disabling RUM
- D. Only alert on RUM

Answer: A — They complement: proactive coverage plus real-user truth.

Dashboards, Metrics, Notebooks & DQL (20 questions)

Q166. Classic dashboards in Dynatrace are composed of:

- A. Tiles bound to metrics, problems, logs, or markdown
- B. Physical widgets
- C. Only text
- D. Spreadsheets

Answer: A — Dashboards arrange tiles visualizing selected data.

Q167. The newer platform 'Dashboards' app additionally supports:

- A. DQL-powered tiles and variables
- B. Only static images
- C. No filters
- D. VHS export

Answer: A — New dashboards can be driven directly by DQL queries and variables.

Q168. Notebooks are best described as:

- A. Interactive documents combining DQL queries, visualizations, and text for analysis
- B. Paper pads
- C. Only PDFs
- D. Login pages

Answer: A — Notebooks support exploratory, shareable data analysis.

Q169. Which DQL command starts a query on log data?

- A. fetch logs
- B. select * from logs
- C. GET /logs
- D. tail logs

Answer: A — DQL pipelines begin with fetch, e.g., fetch logs.

Q170. In DQL, commands are chained using:

- A. The pipe symbol |
- B. Semicolons
- C. Commas
- D. Arrows

Answer: A — DQL is pipeline-based; each stage transforms the previous result.

Q171. Which DQL command filters records?

- A. filter
- B. where-only
- C. grep
- D. cut

Answer: A — filter applies boolean conditions to records.

Q172. To aggregate in DQL (e.g., count by level) you use:

- A. summarize
- B. group only
- C. sum()
- D. reduceAll

Answer: A — summarize computes aggregations, optionally by fields.

Q173. Which DQL command sorts results?

- A. sort
- B. orderby only
- C. arrange
- D. rank

Answer: A — sort orders records by one or more fields.

Q174. Metric keys in Dynatrace look like:

- A. builtin:host.cpu.usage
- B. CPU-USAGE.EXE
- C. metric#1
- D. hostcpu

Answer: A — Metric keys use namespaced dotted notation, e.g., builtin:host.cpu.usage.

Q175. The Data Explorer (classic) is used to:

- A. Chart and analyze metrics with splits and filters
- B. Write Java
- C. Manage users
- D. Install agents

Answer: A — Data Explorer builds metric visualizations for dashboards.

Q176. Which tile type shows open problems on a dashboard?

- A. Problems tile
- B. Disk tile
- C. Font tile
- D. Billing tile

Answer: A — Problems tiles summarize current problem status.

Q177. Dashboard filters/management zone selection affect:

- A. The scope of data shown in tiles
- B. The cluster version
- C. Agent installs
- D. Nothing

Answer: A — Zone/tag/timeframe filters scope every compatible tile.

Q178. Timeframe selection on a dashboard changes:

- A. The evaluated time window for tiles
- B. License terms
- C. Host clocks
- D. Tag values

Answer: A — The global timeframe applies to time-series visualizations.

Q179. Which statement about sharing dashboards is TRUE?

- A. Dashboards can be shared with users/groups or published, subject to permissions
- B. Sharing is impossible
- C. Only via screenshots
- D. Requires root

Answer: A — Dashboards support sharing/publishing with access control.

Q180. SLOs (Service-Level Objectives) in Dynatrace are defined using:

- A. An SLI metric/DQL, a target, and an evaluation window
- B. Random goals
- C. Only Apdex
- D. Manual spreadsheets

Answer: A — SLOs combine an indicator, target percentage, and timeframe, with error-budget tracking.

Q181. Error budget in SLO terms is:

- A. The allowed amount of failure remaining before breaching the SLO target
- B. Money for bugs
- C. CPU quota
- D. A tag

Answer: A — Error budget = allowed unreliability given the target.

Q182. Which DQL function/limit keeps only the first N records?

- A. limit N
- B. top-only
- C. head is the only way
- D. first N rows()

Answer: A — limit truncates the result set to N records.

Q183. fetch metrics vs fetch logs vs fetch events differ by:

- A. The data bucket/type being queried from Grail
- B. Nothing
- C. Font size
- D. Region only

Answer: A — The fetch source selects which data type is queried.

Q184. A markdown tile is useful for:

- A. Adding documentation/links/context to dashboards
- B. Running agents
- C. Creating tokens
- D. Encrypting

Answer: A — Markdown tiles annotate dashboards with text and links.

Q185. Which app helps you explore distributed traces in the new platform?

- A. Distributed Tracing app (Traces)
- B. Fonts app
- C. Billing app
- D. BIOS app

Answer: A — The tracing app analyzes spans/PurePaths in the new UI.

Alerting, Settings, Access & Administration (15 questions)

Q186. Problem notification integrations are configured under:

- A. Settings > Integration > Problem notifications
- B. Smartscape
- C. The host page
- D. BIOS

Answer: A — Notification channels are set up in settings and linked to alerting profiles.

Q187. An alerting profile with a 30-minute delay for slowdown problems means:

- A. Notifications for matching slowdown problems are sent only if the problem persists 30 minutes
- B. Problems close early
- C. Davis waits to detect
- D. Emails arrive faster

Answer: A — Profile delays gate notifications until a problem has lasted the set duration.

Q188. Which access mechanism grants a CI pipeline rights to call Dynatrace APIs?

- A. API token (or OAuth client) with appropriate scopes
- B. Admin password in code
- C. No auth
- D. SSH key to cluster

Answer: A — Scoped tokens/OAuth clients are the supported programmatic auth methods.

Q189. Principle of least privilege for tokens means:

- A. Grant only the scopes needed for the task
- B. Grant all scopes always
- C. Share one token globally
- D. Never expire tokens

Answer: A — Minimal scopes reduce blast radius if a token leaks.

Q190. User permissions in Dynatrace SaaS are commonly managed via:

- A. Groups and policies (IAM), optionally bound to management zones
- B. Editing config files on hosts
- C. Browser cookies
- D. Cannot be managed

Answer: A — Group-based policies control access, optionally scoped by zones.

Q191. SSO integration (e.g., SAML) allows:

- A. Users to authenticate via corporate identity providers
- B. Agents to skip tokens
- C. Free licenses
- D. Faster CPUs

Answer: A — SAML/SSO centralizes authentication with the corporate IdP.

Q192. Audit logs in Dynatrace record:

- A. Configuration changes and important user/system actions
- B. User keystrokes everywhere
- C. Only logins
- D. Nothing

Answer: A — Audit logs track who changed what and when.

Q193. Settings 2.0 schemas allow configuration to be:

- A. Managed consistently via UI and API at environment/entity scopes
- B. Only clicked
- C. Printed
- D. Randomized

Answer: A — Settings 2.0 exposes uniform, API-manageable configuration objects.

Q194. Monitoring as Code (e.g., Monaco/Terraform) is used to:

- A. Version and automate Dynatrace configuration deployments
- B. Compile Java
- C. Replace OneAgent
- D. Draw diagrams

Answer: A — Config-as-code tools apply Dynatrace configuration reproducibly across environments.

Q195. Data retention in Dynatrace depends on:

- A. Data type and bucket/retention configuration (e.g., Grail bucket retention)
- B. Moon phase
- C. Host color
- D. Tag count

Answer: A — Retention varies by data type and configured bucket policies.

Q196. A maintenance window scoped by tag will:

- A. Suppress alerting only for entities matching the tag during the window
- B. Stop all monitoring globally
- C. Delete problems
- D. Change tags

Answer: A — Scoping limits suppression to the intended entities.

Q197. Which is the recommended way to give a new team visibility into ONLY their app?

- A. Create a management zone for their entities and bind their group to it
- B. Give admin to everyone
- C. Email screenshots
- D. Separate cluster per team always

Answer: A — Zone-bound groups deliver least-privilege visibility.

Q198. Token best practices include:

- A. Expiration dates, secure storage, rotation, and minimal scopes
- B. Posting in chat
- C. Hardcoding in public repos
- D. Reusing one token forever

Answer: A — Lifecycle management of tokens is a security essential.

Q199. The Dynatrace API can be used to:

- A. Automate entity queries, config, metric ingest/export, and problem handling
- B. Only download PDFs
- C. Change laptop wallpaper
- D. Nothing

Answer: A — Comprehensive APIs support automation across the platform.

Q200. Which statement about environment vs. account level is TRUE?

- A. Account level manages users/licenses across environments; environment level holds monitoring config/data
- B. They are identical
- C. Environment manages billing only
- D. Account stores PurePaths

Answer: A — Account management spans environments; each environment holds its own monitoring configuration and data.